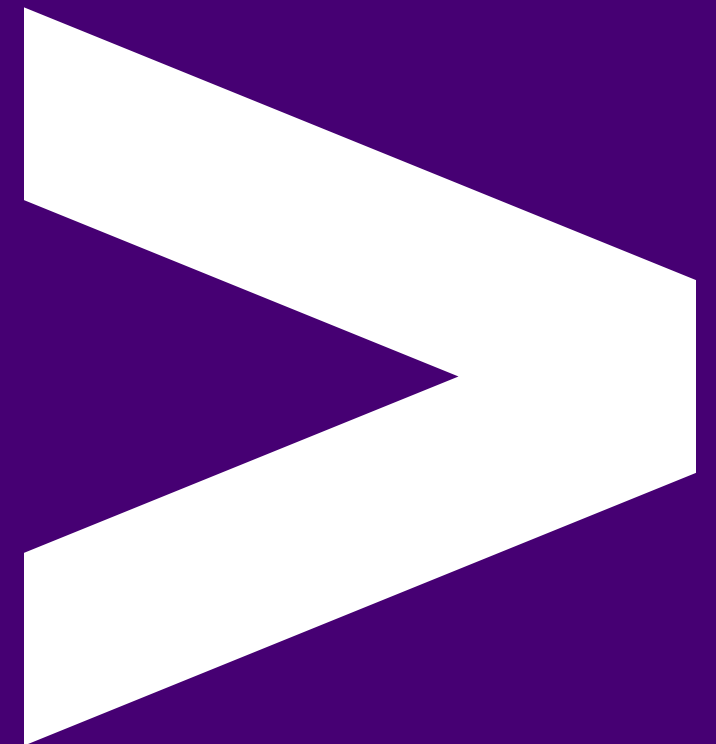








# AWS 06 - Lambda with IaC



## AWS sessions list

- AWS 01 AWS + Cloud Intro  1.5hrs
- AWS 02 AWS CLI Setup  1.5hrs
- AWS 03 S3 Storage (Console)  1.5hrs
- AWS 04 CloudFormation Intro + S3 Storage (IaC)  1.5hrs
- AWS 05 Lambda Intro  1.5hrs
- AWS 06 Lambda (IaC)  1.5hrs
- AWS 07 Redshift (IaC) 1.5hrs
- AWS 08 EC2 (IaC) + Grafana setup 1.5hrs

# Overview

- Lambda as compute in AWS
- Packaging with CloudFormation
- Deployment with CloudFormation
- Event Triggers for Lambdas (with S3)

# Learning Objectives

- How to add a lambda function using CloudFormation
- How to add a trigger between the S3 bucket and the Lambda
- How to process the incoming event in your Lambda function

# Lambda & CloudFormation

This...

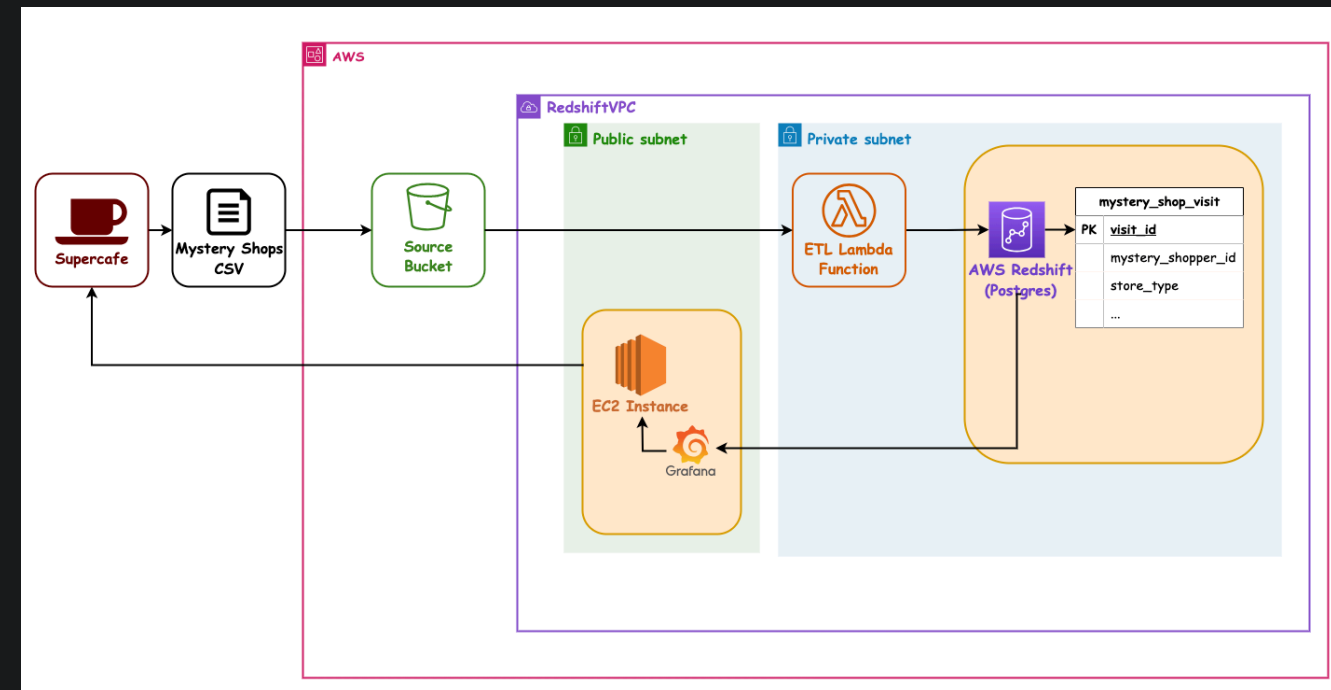


...and this



# Proposed Pipeline Architecture

Let's revisit our Mystery Shopper target setup:



## Our next user story (same as last session)

As a SuperCafe senior manager

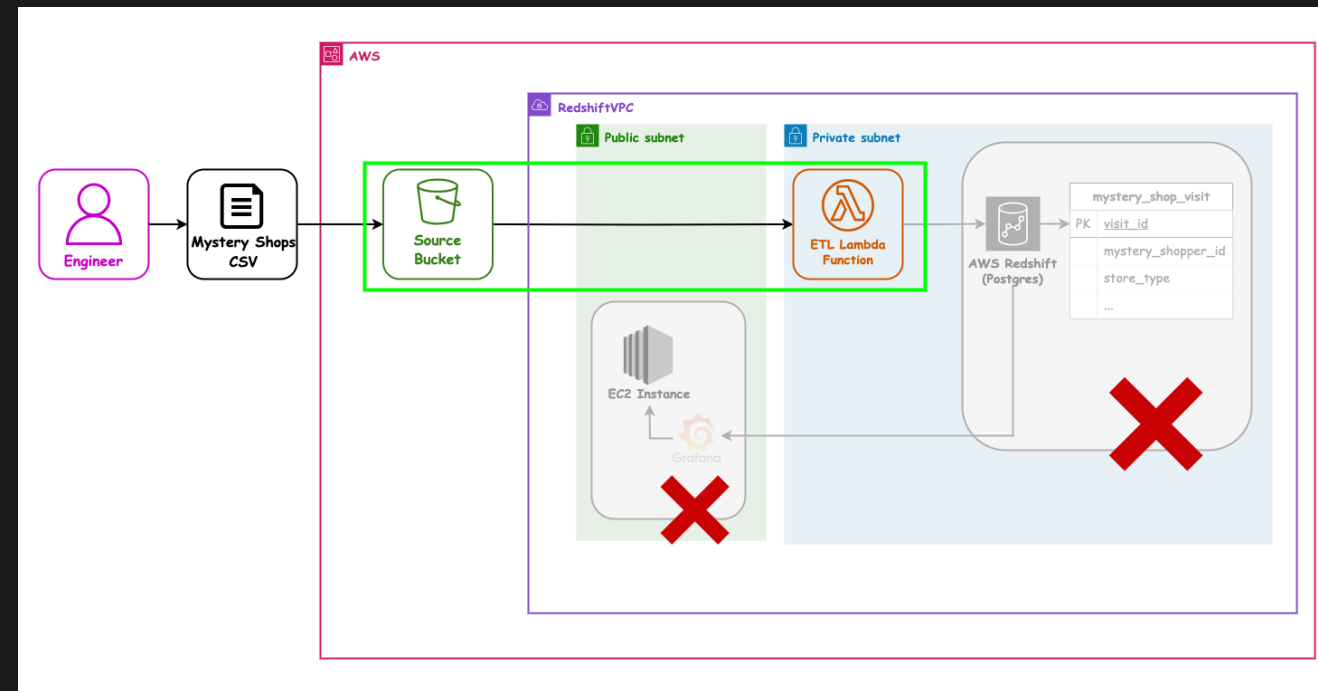
I want the Mystery Shopper data processed automatically

So that the data can be analysed

And the pipeline can run daily

# Our next user story - Architecture

Now that we know Lambda a bit, we will deploy a Lambda "properly" using IaC:

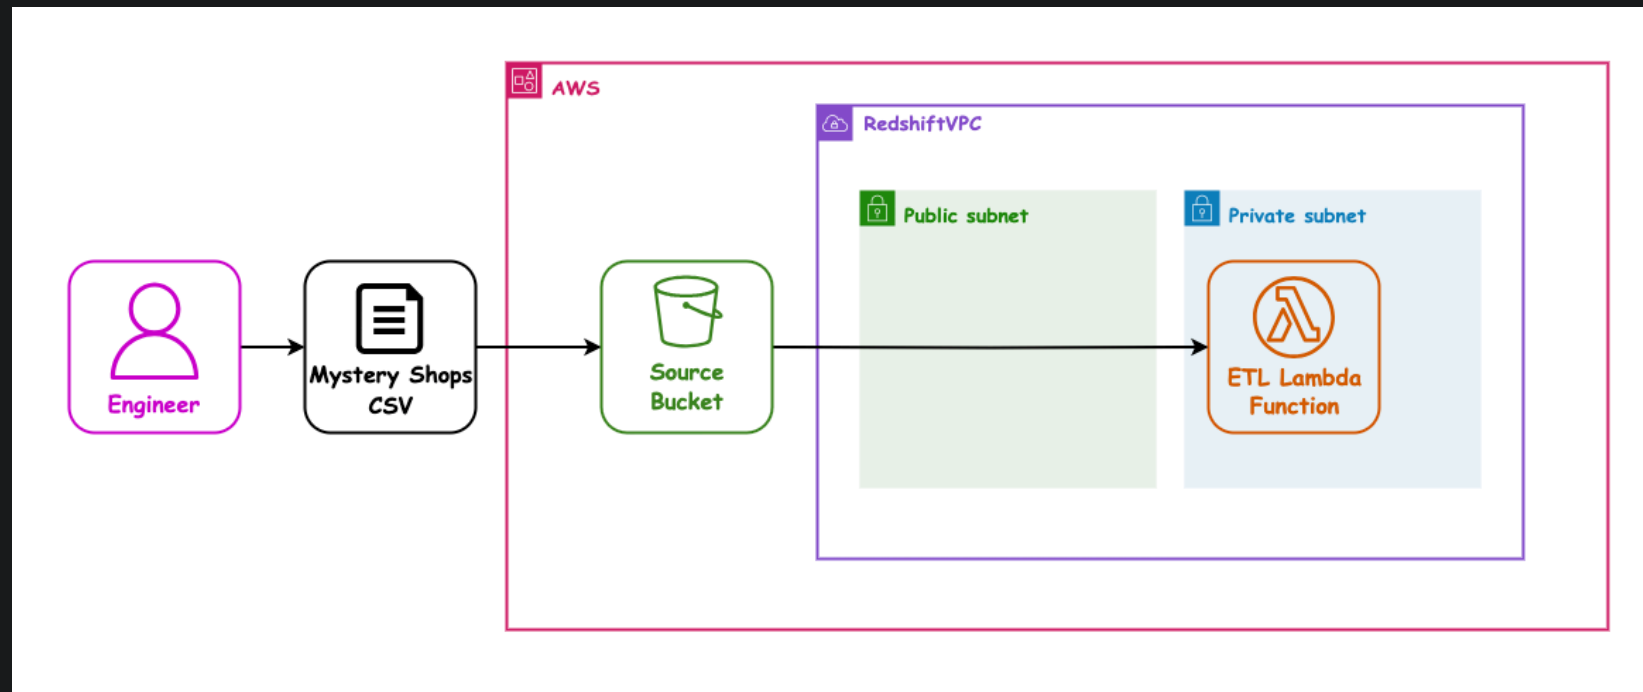


This session - we will set it up with IaC.



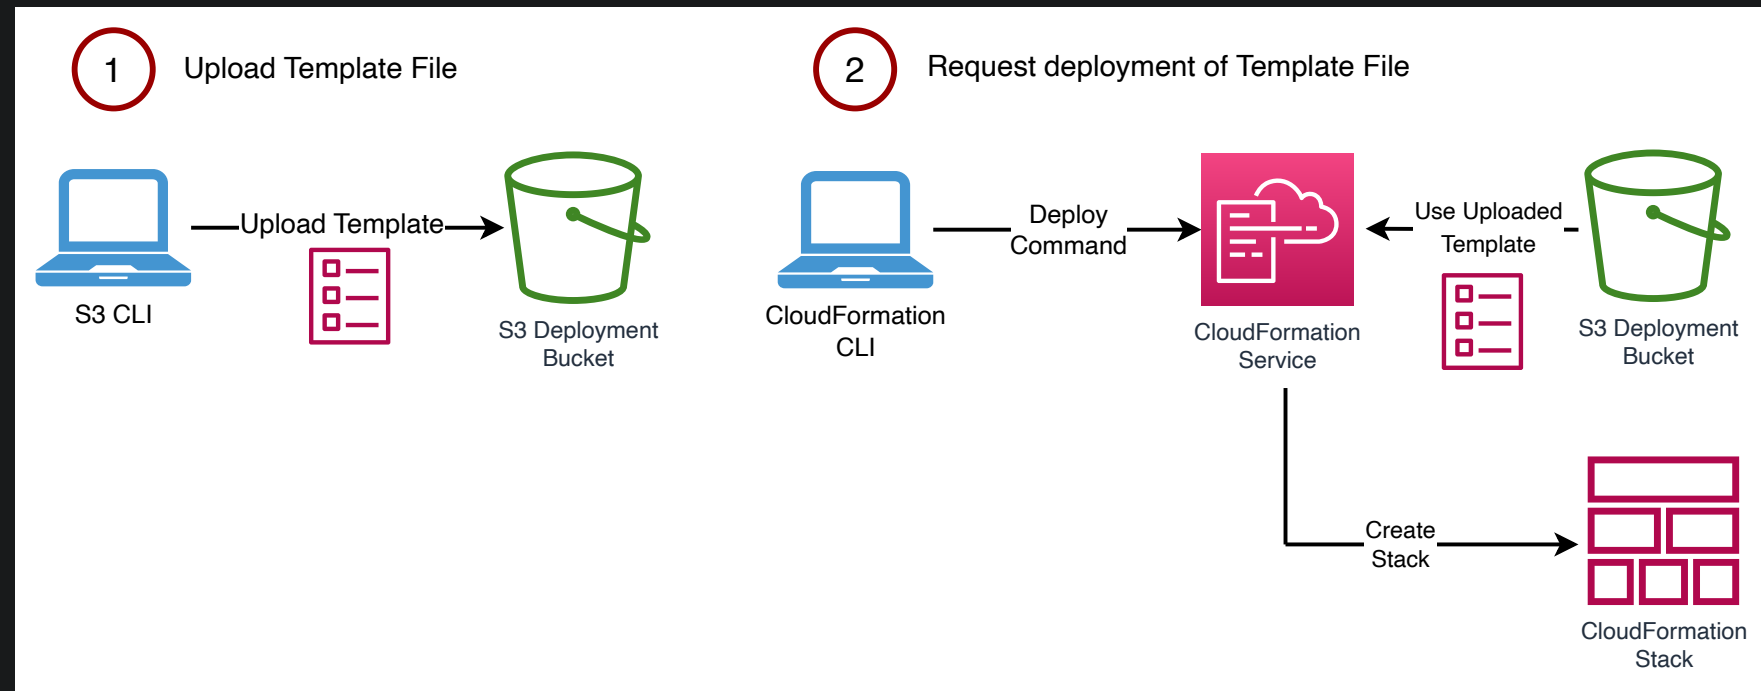
# Our next user story - Architecture

These are the pieces we will need this session:



# Cloudformation Deployment

This is a complex process with a few stages:



*Some notes on the next slides.*

# Deployment Bucket

Our lambda code can get very big, especially with added dependencies. It is common practice to do the following when deploying lambdas from IaC:

- Install any dependencies locally, into the same folder as our python code
- Zip up the Lambda code folder, including the above dependencies
- Upload the Zip to a "Deployment Bucket"
- The Lambda is then deployed from the Zip in the Deployment bucket into the Lambda service

# Packaging vs Deployment

Packaging is the act of getting CloudFormation to:

- Bundle our Lambda code into Zip files,
- Upload the zip files somewhere ready to use later (S3)
- Update our template YAML files to point to the Zip, so that CF knows what to do in AWS

We can see this in the [./handouts/deploy.sh](#) file.

# Our Data Bucket

In our previous session(s) we set up a *data* bucket to put our CSV files in.

This is separate to the *Deployment* Bucket we need to put our zips of lambda code into.

# Demo - Starting point

We need to start from the partially complete file [../handouts/etl-stack.yml](#).  
This continues from our previous sessions.

## Demo - The provided Lambda

The Lambda code to process the CSV has been written for us by our team mates - we can see this in file [./handouts/src/mystery\\_shop\\_etl\\_lambda.py](#).

*This session is about using CloudFormation for Lambda, as distinct from writing the workings of a python lambda ourselves.*

# Demo - The provided data

A sample [mystery\\_shops\\_2024-03.csv](#) file is also provided.

```
mystery_shops_2024-03.csv ×
data-academy-pipeline-example > data > mystery_shops_2024-03.csv > data
1 StoreID,StoreName,MysteryShopperID,MysteryShopperName,StoreType,NumberOfStoreEmployees,VisitDate,Star
  rtTime,EndTime,OverallScore
2 1487,Leeds,22,Rory Gilmore,Free Standing,"12",13/03/2024,10:05,10:30,3
3 1456,London,48,Lane Kim,Mall,"6",21/03/2024,13:45,15:00,2
4 1403,Manchester,32,Luke Danes,Mall,"7",22/03/2024,10:30,10:45,5
5 1482,Newcastle,32,Luke Danes,Mall,"7",22/03/2024,14:02,14:17,5
6 1499,Edinburgh,22,Rory Gilmore,Kiosk,"4",19/03/2024,15:58,16:20,4
7
```



# Code along - Parameter

Add a parameter for Network Stack Name, so we know where to put the lambda (so that in a later session it can talk to RedShift).

- In the `Parameters` section
- With logical name `NetworkStackName`
  - With a `Type` of `String`
  - A `Default` value of `project-networking`
  - And a helpful `Description`

# Code along - lambda

Add a Lambda with a dynamic name (from `YourName`), so all our lambdas are unique.

- In the `Resources` section
- With a logical name like `EtlLambdaFunction`
- And a specific `Type` of `AWS::Lambda::Function`
- And many `Properties...`

*See next slide for more.*

# AWS Lambda properties

There are many that we can set, we need at least the following:

- A unique and dynamic `FunctionName`, using `YourName`
- An up to date `Runtime`, `python3.12`
- A `Handler` to specify the file name and function name to run
- The `Code` setting, to specify which folder our source code is in e.g. `./src`
- A `Role`, to assume for security so we are allowed to talk to RedShift and the S3 bucket
- A `Timeout` value in seconds e.g. `30`, high enough for our E-T-L to run but not time out
- A `VpcConfig`, to put our lambda in the same networking as RedShift so it can see the DB
- A `Tag` with value `Name` to further identify our lambda

# Code along - Wake the Lambda

Add a Notification Configuration to the CSV data bucket, so that files arriving there wake up the lambda.

We need to extend the `ShopperRawDataBucket` configuration `Properties`, like so:

- Add a new `NotificationConfiguration` property
- With a child property of `LambdaConfigurations`
  - This has a child list of Event & Function tuples
  - Add an `Event` of type `s3:ObjectCreated:*`
  - With a `Function` (lambda) reference to `!GetAtt EtlLambdaFunction.Arn`

## S3 as a source

As mentioned in the previous session [aws-03-console-s3](#), S3 can be a source of our data:

- We've set up S3 to send a notification or "event" to our Lambda to wake it
- This event will tell us the *bucket* and *file name* (but not the content / payload)
- If those systems fail to respond, some of them will receive a retry - for example, if Lambdas are *throttled*, S3 will retry the event for up to 6 hours

# Code along - Dependencies

We will tell CloudFormation that the Bucket depends on the permissions and the lambda.

```
ShopperRawDataBucket:
  ...
  DependsOn:
    - ShopperRawDataBucketPermission
    - EtllambdaFunction
```

Most of the time, CloudFormation will work these out for it's self. However we have found in this stack, the build order is more reliable with this hint added.

# Code along - Because Security

Add a Source Bucket Permission, so the Lambda is allowed to read from in the bucket when it is invoked.

- We need a new Resource called `ShopperRawDataBucketPermission`
- With a Type of `AWS::Lambda::Permission`
- The Properties of it are
  - An Action, which is `lambda:InvokeFunction`, for when the lambda is activated
  - The `FunctionName`, by reference to our lambda, e.g. `!Ref EtlLambdaFunction`
  - For the specific `Principal` that is `s3.amazonaws.com`
  - Allowing the `SourceArn` by name so `!Sub 'arn:aws:s3:::${YourName}-shopper-raw-data'`

# Code along - Log into AWS

Make sure you are logged into AWS in your terminal

```
aws-azure-login --profile sot-academy
```

- Windows users may need to use Powershell



# Demo - the Deploy script - 5 mins

The deploy script [./handouts/deploy.sh](#) is done for you, so that it will reliably work.  
Instructor to show the file.

It does the following:

- Collect your `aws-profile` and `your-name` from the command line
- Deploy a stack called `your-name-shopper-deployment-bucket`
- Install the Lambda's dependencies in the `src` folder
- Package the `your-name-shopper-etl-pipeline` stack with Lambda Zip in S3
- Deploy a stack called `your-name-shopper-etl-pipeline`

# Code along - Deployment

Let's all deploy our stacks. This may take some time!

- Windows users may need to do this in GitBash
- **YourName** should be entered **lower-case-with-dashes**, as it will be used in the S3 Bucket names

Run the [./handouts/deploy.sh](#) script like this:

```
cd handouts
./deploy.sh <aws-profile> <your-name>
# e.g.
./deploy.sh sot-academy rory-gilmore
```

# Code along - Trigger the lambda

Upload the sample CSV file [mystery\\_shops\\_2024-03.csv](#) into your data bucket.  
This should trigger your lambda.

# Code along - Check the logs

Find your Log Group in CloudWatch and check the latest Log Stream.  
...Do you see some nice useful logs?

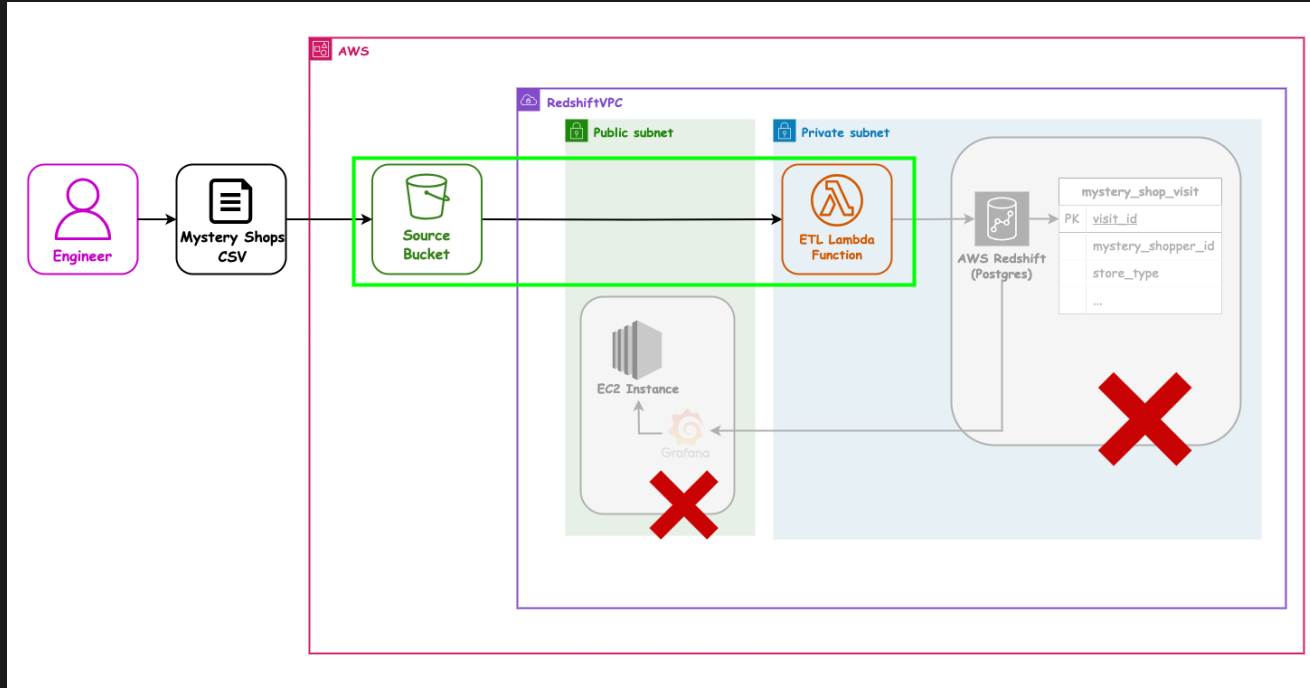
# The results

In the `./solutions` folder there is a completed `etl-stack.yml` with extra comments, as a refresher of what we have assembled.

*This is provided so that after the session you can cross-reference what we put together with the slides.*

# Our next user story - Architecture

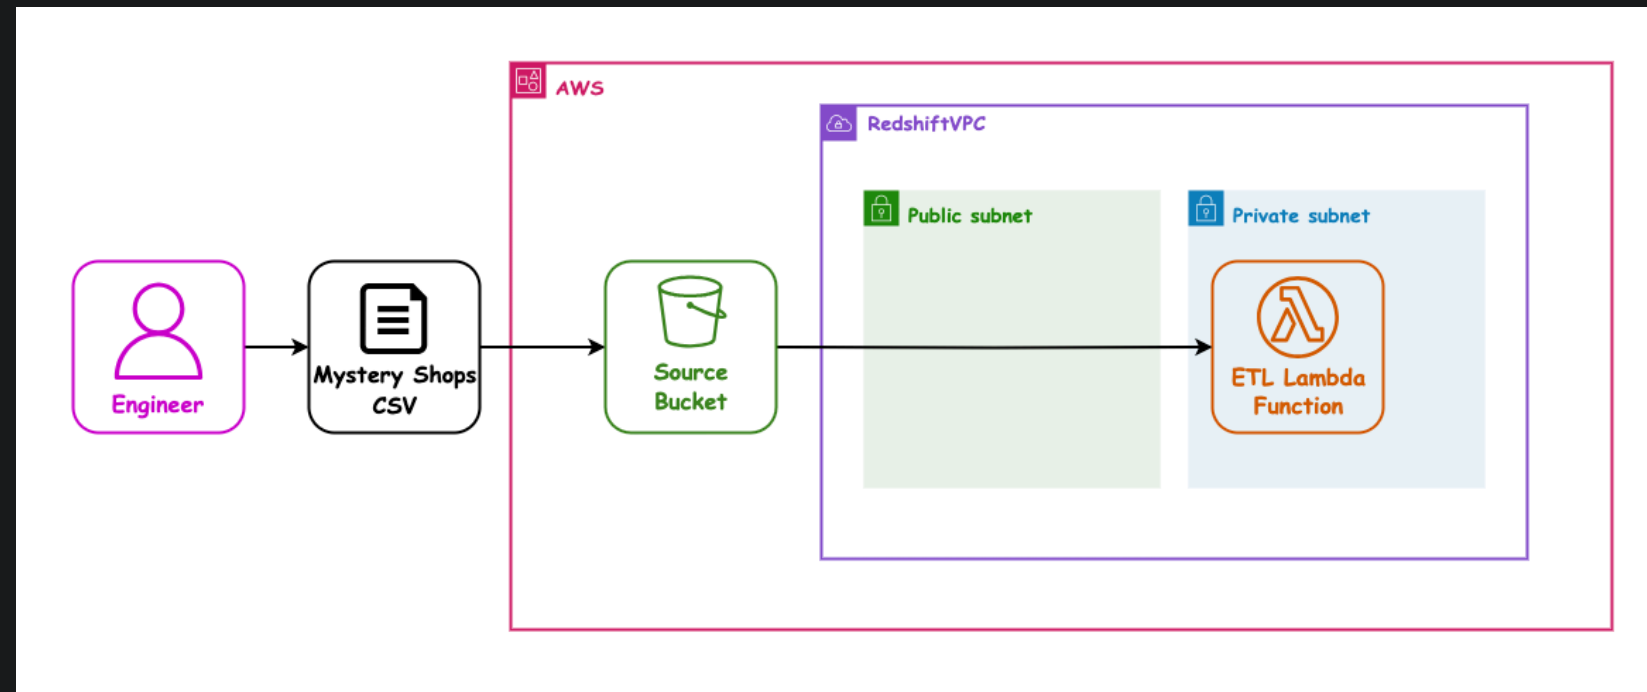
## This is what we just did with IaC:



This session - we set it up with IaC instead of manually.

# Our next user story - Architecture

These are the bits we used:



# Terms and Definitions - recap

- Lambda
- Event Driven
- Event Trigger
- Handler function
- S3
- Package
- Deployment



# Overview - recap

- Lambda as compute in AWS
- Packaging with CloudFormation
- Deployment with CloudFormation
- Event Triggers for Lambdas (with S3)

## Learning Objectives - recap

- How to add a lambda function using CloudFormation
- How to add a trigger between the S3 bucket and the Lambda
- How to process the incoming event in your lambda function

## Emoji Check:

On a high level, do you think you understand the main concepts of this session? Say so if not!

1. 🥲 Haven't a clue, please help!
2. 😞 I'm starting to get it but need to go over some of it please
3. 😐 Ok. With a bit of help and practice, yes
4. 😊 Yes, with team collaboration could try it
5. 😄 Yes, enough to start working on it collaboratively