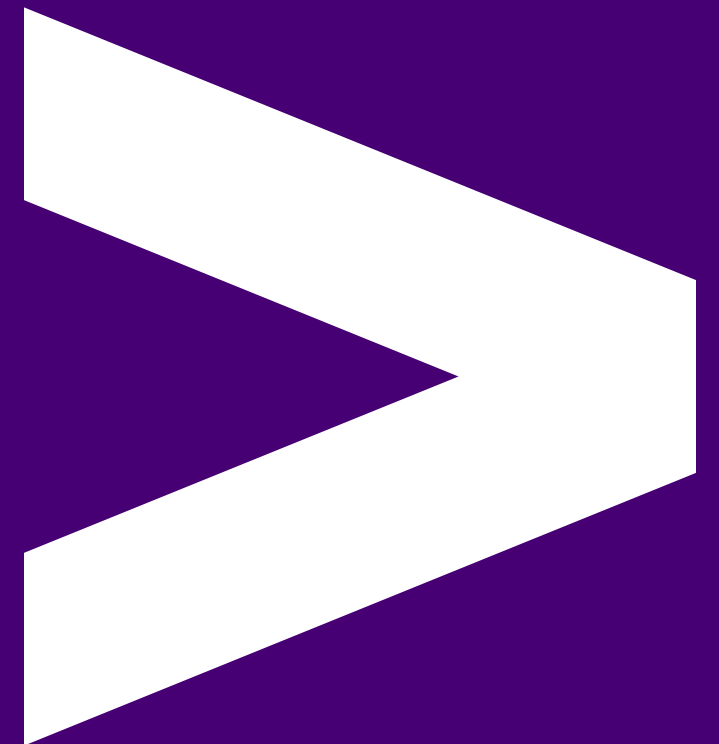









AWS 07 - Introduction to Amazon Redshift



AWS sessions list

- AWS 01 AWS + Cloud Intro  1.5hrs
- AWS 02 AWS CLI Setup  1.5hrs
- AWS 03 S3 Storage (Console)  1.5hrs
- AWS 04 CloudFormation Intro + S3 Storage (IaC)  1.5hrs
- AWS 05 Lambda Intro  1.5hrs
- AWS 06 Lambda (IaC)  1.5hrs
- AWS 07 Redshift (IaC)  1.5hrs
- AWS 08 EC2 (IaC) + Grafana setup 1.5hrs

Overview

- Before Data Warehouses
- What is a Data Warehouse?
- Why use a Data Warehouse?
- Redshift as a Data Warehouse
- Connecting from a lambda using Python

Learning Objectives

- What Redshift is and can do for us
- How to connect to it using Python
- Using Parameter Store to get settings
- Never put sensitive info like passwords in code or environment variables!

Amazon Redshift



Before Redshift - Traditional Data Warehouses

- Time consuming to pull data from many sources using traditional architecture
- Costly - hardware, setup, electricity, security, estate
- Maintenance costs often outweighed the benefits (upgrading systems due to more data being added)
- Performance issues
- Auto-scaling is not an easy concept

Discussion (5 mins)



What is a data warehouse?

Why would we use it?

Data Warehousing

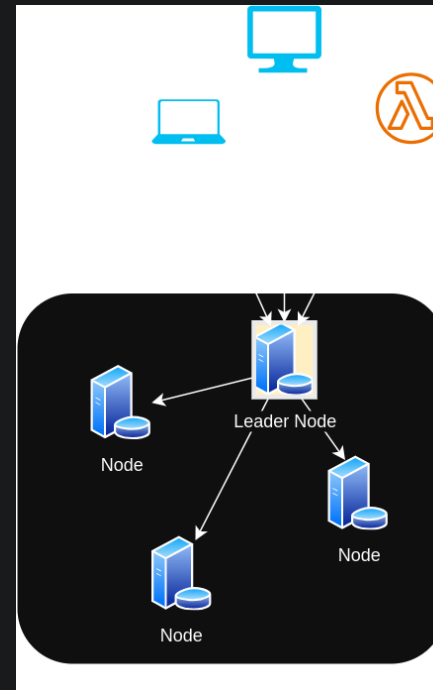
- It intends to be where you keep your best data - the source of truth for your business
- It has been cleaned and transformed in preparation for storage
- Highly structured data that is ready for analysis

Redshift

- Massively **parallel, column-oriented** database
- Simple and cost-effective to analyse your data
- Manages, monitors and scales your system
- Up to 10x better performance than traditional
- Collection of compute resources which are called nodes
- These nodes, when organised into groups, become **clusters**
- Each cluster runs a Redshift engine which contains one or more DBs

Architecture

RedShift operates a complex model under the hood:




Architecture

Sample data



Architecture

Sharding



Clusters

- A cluster has a leader node with one or more compute nodes
- The leader node receives queries from client applications (BI, analytical software etc.)

Compute Nodes

- Compute resources which execute a query plan
- Transmits data among themselves to solve queries
- Nodes are further divided into (node) slices
- Each node slice receives an allocation of memory and performs operations in parallel

Node types

- When you launch a (non-free tier) cluster, you need to specify the node types
- There are several types of node you can pick, including:
 - RA3
 - DC2 (Dense Compute)
 - DS2 (Dense Storage - *Legacy*)

Choosing the right nodes:


- **Data Quantity:** Be aware of the amount of data you want to import into your Redshift cluster
- **Complexity of queries:** Different nodes support queries with differing levels of complexity
- **Downstream systems:** What uses the results of the queries? How important is query speed?

Columnar Data Storage

- The data is still represented with rows and columns as normal
- However, the data is physically stored by column, instead of rows
- Because the data stored is the same type, you can achieve better data compression
- Number of I/O operations decreases
- Also means you can query/perform data analysis on similar types of data far quicker than row storage

Proposed Pipeline Architecture

Let's revisit our Mystery Shopper target setup:



Our next user story

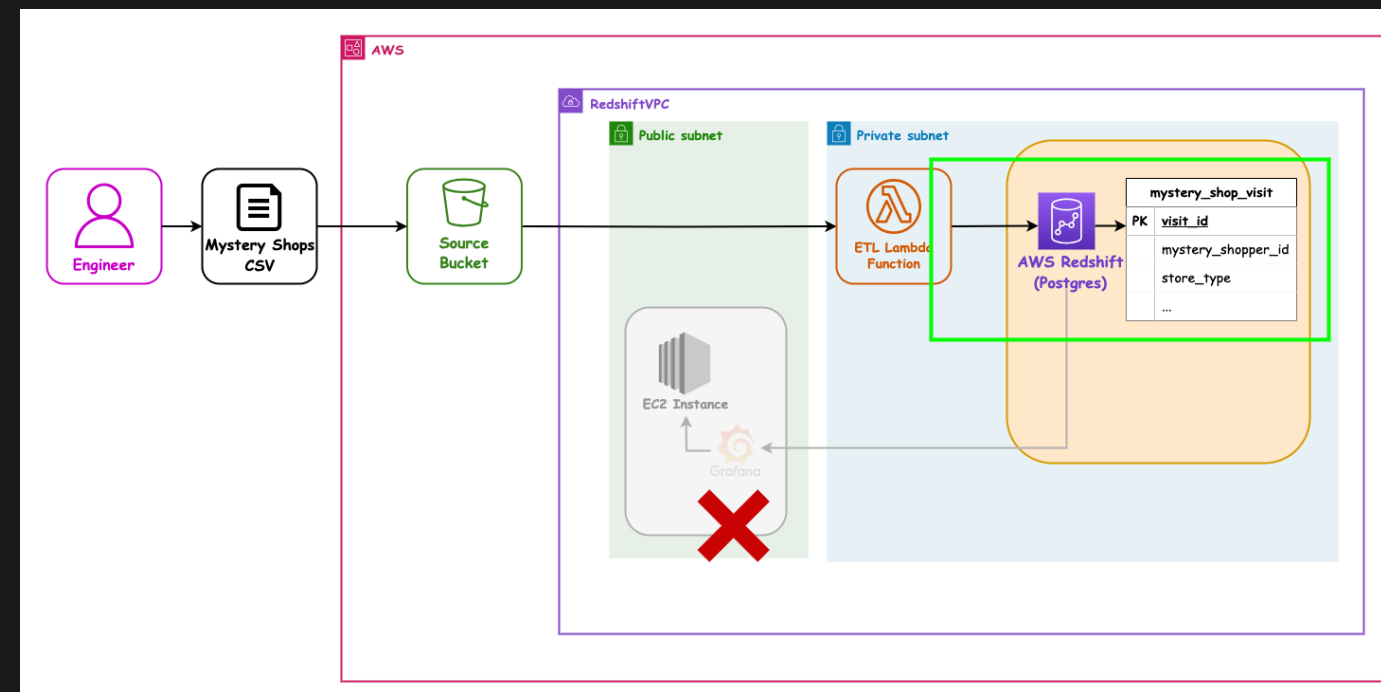
As a SuperCafe senior manager

I want the Mystery Shopper data to be analysed

So that we can award the friendliest most helpful store a prize each month

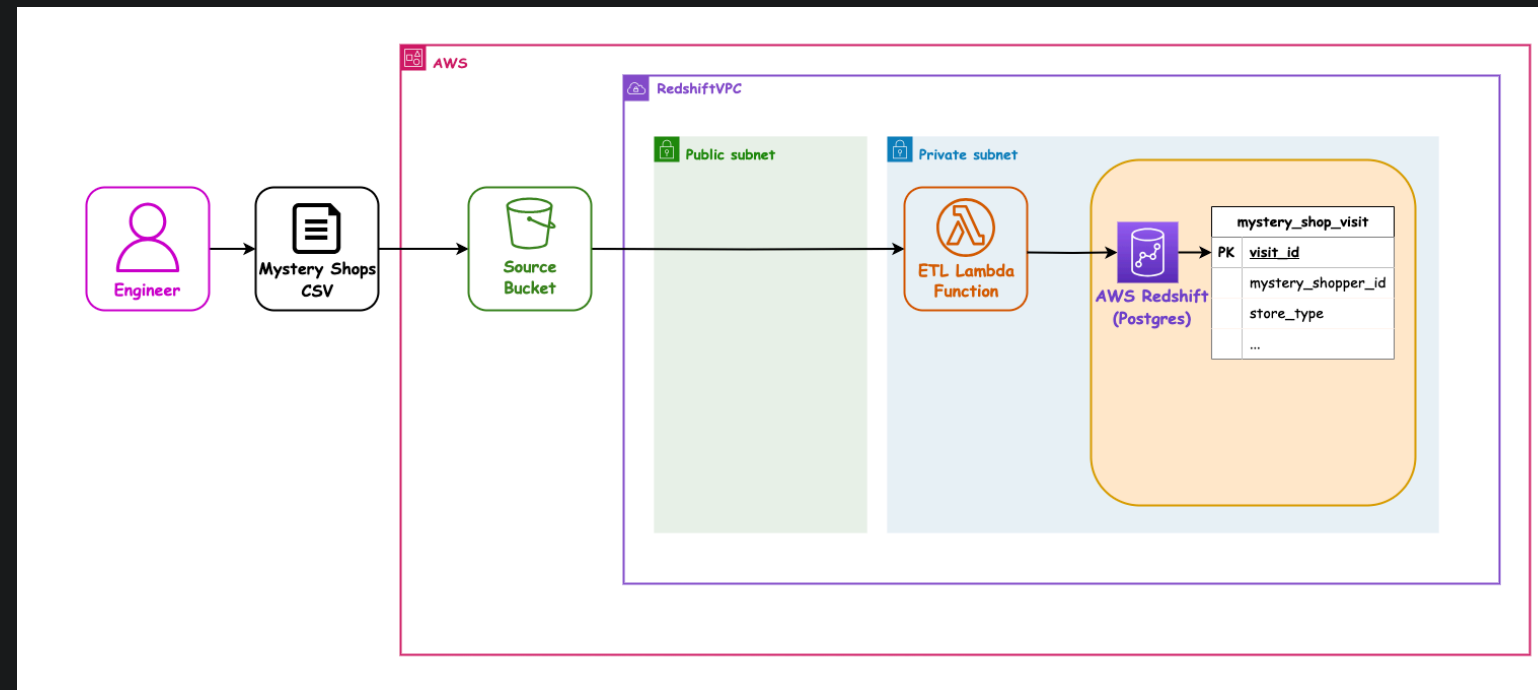
Our next user story - Architecture

We want to check our data is going into the database, so that we can query it. We don't need the Grafana VM yet:



Our next user story - Architecture

So for now we only need this chain of parts to be working:



Wait, no new CF templating?

This session does not involve new CloudFormation *per-se* - rather, using tools already in our toolbox, we connect to RedShift, add data, and query it.

The ETL Lambda python code provided in this session has been extended to talk to RedShift for us.

Exercise time

Now we've learned a little bit about Amazon Redshift, let's implement loading some data from the Mystery Shopper pipeline into a database hosted on a Redshift cluster.

Breakouts will be in each Project Team.

Demo - new lambda code

We need to deploy the updated lambda code.

Instructor to show the new lambda code.

Especially important is the use of Parameter Store to get the connection details.

- See [./handouts/src/mystery_shop_etl_lambda.py](#).
- And [./handouts/src/utils/db_utils.py](#).
- And [./handouts/src/utils/sql_utils.py](#).

Exercise Prep: Your team name

Your team name will be provided by the instructors. It will have dashes in it (–) due to limitations in AWS (esp. Bucket names).

E.g if your team name is "La Vida Mocha", your logical team name will be `la-vida-mocha`.

You will need this!



Note the dashes!

Demo - Parameter Store settings

The RedShift connection details will be in a Parameter Store param called `<team_name>_redshift_settings`,
e.g. `la_vida_mocha_redshift_settings`, as Redshift likes underscores (_).

Instructor to show each team their Parameter Store settings,
e.g. `la_vida_mocha_redshift_settings`.

Demo - the Deploy script - 5 mins

The deploy script [./handouts/deploy.sh](#) is done for you, so that it will reliably work.
Instructor to show the file.

It does the following:

- Collect your `aws-profile` and `team-name` from the command line
- Deploy a stack called `team-name-shopper-deployment-bucket`
- Install the Lambda's dependencies in the `src` folder
- Package the `team-name-shopper-etl-pipeline` stack with Lambda Zip in S3
- Deploy a stack called `team-name-shopper-etl-pipeline`

Exercise - Step 01 - Deploying

15 mins breakout

Only ONE person should do this with the rest of the team watching... ...have a think about why!

- Follow [Exercise Step 01 - Deploying](#)
- Deploy the new lambda code using your team name e.g. `la-vida-mocha` (lower case with dashes)
- Run `./deploy.sh <profile-name> <team-name>` to update the ETL lambda function:
 - e.g. `./deploy.sh sot-academy la-vida-mocha`

Discussion - Step 01 - Deploying

How did the deployment go?

Is everyone happy they understand it?

- Why could only **ONE** person from each Team do this?

Exercise - Step 02 - Trigger lambda

5 mins breakout

- Follow [Exercise Step 02 - Trigger lambda](#)
- Use file [./handouts/data/mystery_shops_2024-03.csv](#)
- You can upload the same file many times, or altered copies of it
- Multiple team members can do this concurrently
 - Why is that?

Discussion - Step 02 - Trigger lambda

How did the triggering go?

Is everyone happy they understand it?

- Why could multiple people do this at the same time?

Demo - Step 03 - RedShift Query Editor

The Instructor will now demo how to connect to the RedShift query editor in the AWS Console, using the Parameter Store data, e.g.
`la_vida_mocha_redshift_settings`.

Exercise - Step 03 - Create SQL Queries

- 20 mins breakout

Follow [Exercise Step 03 - Create SQL Queries](#).

Use the RedShift query editor in the AWS Console.

- You can all do this concurrently (not just one person per team)
- Start with `select * from mystery_shop_visit;` and then see if you can do something more sophisticated

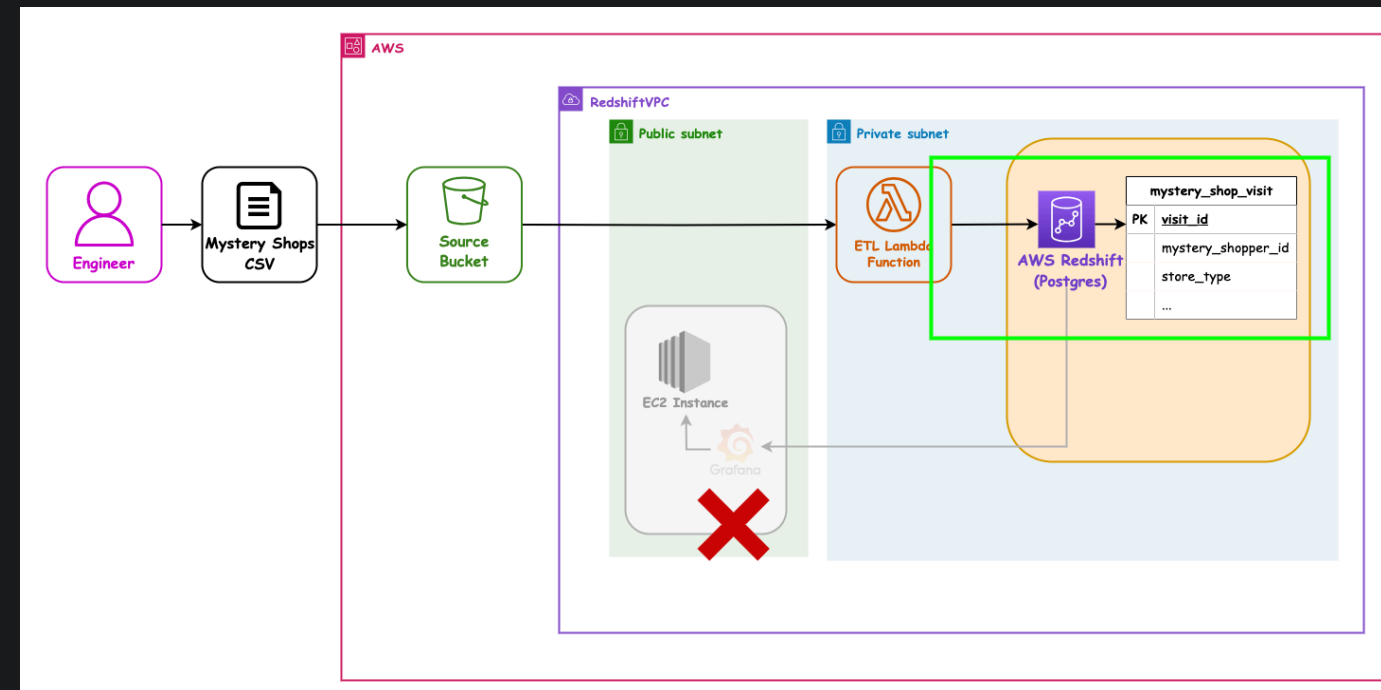
Discussion - Step 03 - Create SQL Queries

How did the SQL-ing go?

Is everyone happy they understand the queries?

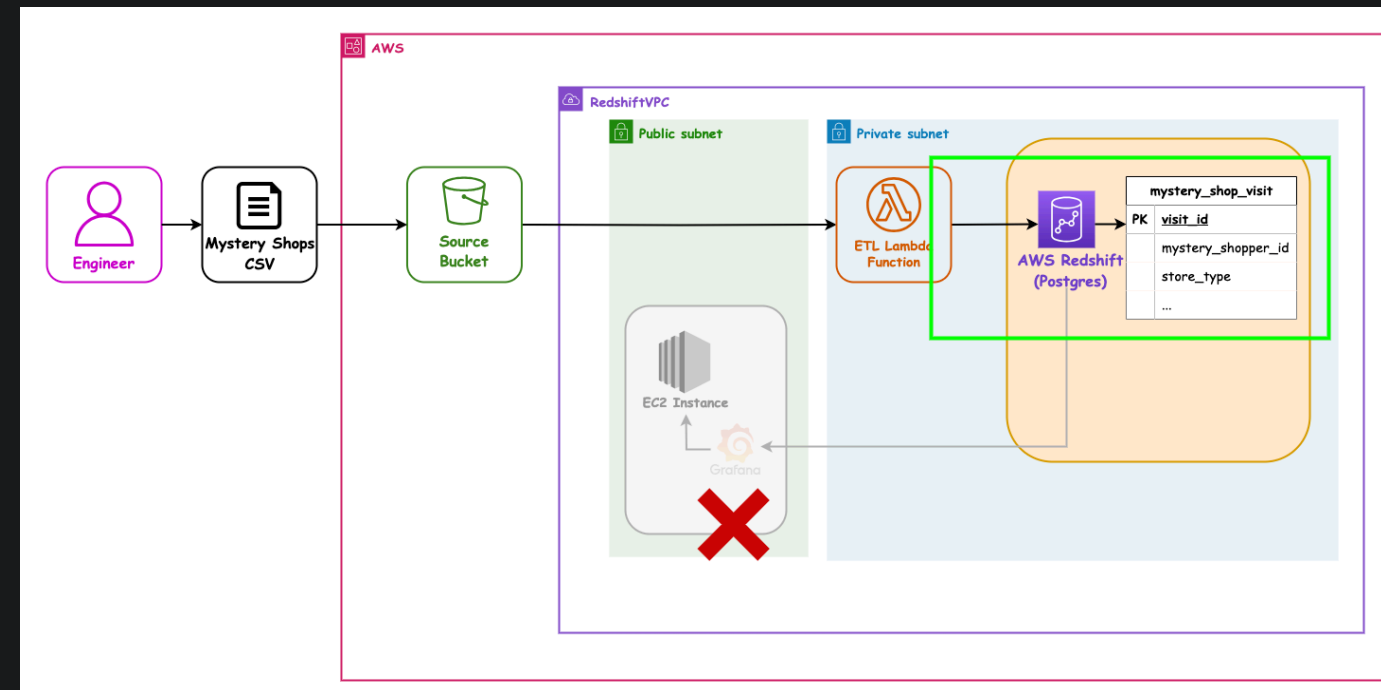
Our next user story - Architecture

This is what the IaC has set up and allowed us to use:



Our next user story - Architecture

We nearly have the whole system set up now:



Overview - recap

- Before Data Warehouses
- Redshift as a Data Warehouse
- Connecting from a lambda using Python

Learning Objectives - recap

- What Redshift is and can do for us
- How to connect to it using Python
- Using Parameter Store to get settings
- Never put sensitive info like passwords in code or environment variables!

Further Reading

- **Redshift:** The AWS Data Warehouse service used to store and analyse large quantities of data
- AWS Docs - <https://aws.amazon.com/redshift/>

Emoji Check:

On a high level, do you think you understand the main concepts of this session? Say so if not!

1. 🥲 Haven't a clue, please help!
2. 😞 I'm starting to get it but need to go over some of it please
3. 😐 Ok. With a bit of help and practice, yes
4. 😊 Yes, with team collaboration could try it
5. 😄 Yes, enough to start working on it collaboratively